

Ternlang: A Full-Stack Balanced Ternary Execution Architecture for Sparse Neural Inference and Ambiguity-Aware Agent Systems

The Autonomous Substrate — The Death of the Bit

1
—

Abstract

We present **Ternlang**, the first complete software stack for balanced ternary computing: a domain-specific language, bytecode compiler, stack-based virtual machine (BET VM), hardware description language backend, distributed actor runtime, and machine learning inference kernels—all unified under a single coherent architecture. The foundational primitive is the *trit* $(t \in \{-1, 0, +1\})$, where the value (0) represents an *active neutral state* rather than absence, enabling three-valued logic that is structurally superior to binary for ambiguity-aware reasoning.

The principal contribution is **TSPARSE_MATMUL**: a first-class VM opcode that elides multiply operations against zero-weighted (“hold”) trit elements, surfacing at the instruction-set level a property that BitNet-style ternary quantization reveals in neural weights. Empirical evaluation on quantized weight matrices of (512×512) elements yields (56.2%) sparsity and a $(2.27\times)$ reduction in multiply operations versus dense execution; a Compressed Sparse Column kernel with Rayon parallelism achieves up to $(122\times)$ speedup at (99%) sparsity—without approximate arithmetic or hardware reconfiguration.

We further define the **Balanced Ternary Execution (BET) ISA**, a formal 2-bit-encoded instruction set with 51 opcodes spanning arithmetic, tensor operations, actor messaging, and control flow; synthesise it to Verilog-2001 with per-cell clock-gating on zero-weight elements; and demonstrate interoperability with existing ternary computing efforts via the `ternlang-compat` bridge crate.

Beyond the execution substrate, we introduce two higher-level contributions. The **Ternary AI Reasoning Toolkit** (Phase 8) provides five structural primitives—iterative deliberation, coalition voting, multi-dimensional action gating, scalar-to-temperature bridging, and hallucination scoring—that make any AI agent’s decision loop architecturally ternary rather than merely prompted. The **MoE-13 Ternary Orchestrator** (Phase 9) implements a Mixture-of-Experts head with dual-key synergistic routing, $(1\{+\}1\{=\}3)$ emergent triad synthesis, a hard-wired safety veto, and a three-tier memory mesh, realising the MoE-13 architecture in 177+ passing tests across the full workspace. All nine crates are published to `crates.io` under open-core licensing. A companion corpus of **28,500+** executable `.tern` programs demonstrates the hold state across domains spanning aerospace, medicine, distributed systems, civic governance, and AI agent design.

Version 2.0.0 of the Ternary Intelligence Stack introduces three further empirical advances. First, **Albert MoE-13**—a 22M-parameter fully ternary Mixture-of-Experts language model—demonstrates autonomous architectural self-evolution via Net2Net safe-copy layer surgery (3L to 5L without restart), achieving an all-time best single-batch cross-entropy of **2.1353** (perplexity ≈ 8.5) trained entirely on a 2013 consumer CPU with no GPU. Second, independently verified AVX2 SIMD hardware benchmarks confirm **1.80 $\times$** throughput at 50% sparsity and **3.0 $\times$** over INT8 at 90% sparsity on a production Axum server. Third, the BET VM has been compiled to WebAssembly and executes live in the browser via TernStudio, proving that the complete ternary execution pipeline—lexer, parser, codegen, VM—runs client-side on standard web hardware without installation. The live MCP surface has grown to **34 free tools**. These results establish ternary computation as a practical, democratising inference substrate rather than a theoretical curiosity.

balanced ternary, ternary logic, sparse inference, BitNet, domain-specific language, virtual machine, actor model, FPGA synthesis, Verilog, mixture-of-experts, AI reasoning, ternary orchestration, autonomous training, self-evolving neural networks, consumer CPU inference, ecocentric AI, WebAssembly

Introduction

The computational substrate underlying modern artificial intelligence is binary. Floating-point arithmetic, two-state memory, and Boolean logic gates have driven five decades of progress—but they introduce a fundamental representational mismatch when modelling systems that are inherently three-valued: *affirmed*, *denied*, and *undecided*.

Clinical diagnosis, legal reasoning, sensor fusion under noise, and multi-agent consensus all require a native neutral state that binary computing forces to encode as a special case—null pointers, NaN, sentinel values, or probabilistic scores collapsed to a threshold. Each encoding is a workaround for an absent primitive.

Balanced ternary provides that primitive. A *trit* ($t \in \{-1, 0, +1\}$) carries three symmetric values. The neutral value (0) is *active*—a deliberate state of hold, not an empty bit pattern. Balanced ternary arithmetic is self-complementing: negation requires no special-case handling. And at the scale of modern neural networks, where BitNet and related work show that ternary-quantized weights preserve accuracy with dramatically reduced computation, the case for a ternary-native execution substrate is both theoretical and empirical.

Despite this, the ternary computing field remains fragmented: hobbyist emulators, academic EDA tools for memristor hardware, isolated Lisp interpreters, and hardware simulators without compiler support. No project provides the full vertical stack.

Ternlang fills this gap. Our contributions are:

1. A *language design* for balanced ternary: three-way exhaustive pattern matching, first-class trit tensors, and an actor model for ternary message passing (§3).
2. The **BET ISA**: a formal 2-bit-encoded instruction set with opcode coverage from arithmetic through tensor operations and distributed agent control (§4).

- 3. **TSPARSE_MATMUL**: a VM opcode that skips zero-weight multiplications at the instruction level, realising the sparsity benefit of ternary quantization without software overhead (§5).
- 4. A **Verilog-2001 hardware backend** with synthesisable sparse matmul array and full BET processor, plus a cycle-accurate RTL simulator in pure Rust (§6).
- 5. A **Ternary AI Reasoning Toolkit**: five structural primitives—deliberation, coalition vote, action gate, scalar temperature, and hallucination score—that make any AI agent’s decision loop architecturally ternary (§9).
- 6. The **MoE-13 Ternary Orchestrator**: dual-key synergistic routing, $\set{1}{+}{1}{=}3$ emergent triad synthesis, a hard safety veto, and three-tier memory mesh, exposed as MCP tool calls (§10).
- 7. **Ecosystem bridges** connecting existing ternary projects to the BET VM as a common runtime (§14).
- 8. A **reference example corpus**: twenty executable .tern programs demonstrating the hold state in aerospace, medicine, distributed systems, hardware pipelines, civic governance, finance, and AI agent design—published open-source alongside all nine workspace crates on crates.io.

Background: Balanced Ternary

Trit arithmetic

A trit $t \in T = \{-1, 0, +1\}$ participates in the following complete operations :

Definition 1 (Ternary addition with carry). For trits $(a, b \in T)$, define $(s = \mathrm{sum}(a,b))$ and $(c = \mathrm{carry}(a,b))$ such that $(a + b = 3c + s)$, where $(s, c \in T)$.

The complete sum/carry table is:

Balanced ternary addition (sum, carry)			
$(a \backslash b)$	(-1)	(0)	$(+1)$
(-1)	$((+1,-1))$	$((-1, 0))$	$((0, 0))$
(0)	$((-1, 0))$	$((0, 0))$	$((+1, 0))$
$(+1)$	$((0, 0))$	$((+1, 0))$	$((-1,+1))$

Definition 2 (Ternary multiplication). $(\mathrm{mul}(a,b) = a \cdot b)$ under integer multiplication, restricted to (T) : $(\mathrm{mul}(a,b) \in T)$ since $(|a|,|b| \leq 1)$.

Definition 3 (Consensus (ternary OR)). $(\mathrm{cons}(a,b) = a)$ if $(a = b)$, else (0) .

Definition 4 (Negation). $(\mathrm{neg}(t) = -t)$. Note $(\mathrm{neg}(0) = 0)$.

The 2-bit BET encoding

Hardware naturally operates in binary. We encode trits as 2-bit pairs:

BET 2-bit trit encoding

Bits	Value	Meaning
0b01	$\backslash(-1\backslash)$	conflict
0b10	$\backslash(+1\backslash)$	truth
0b11	$\backslash(0\backslash)$	hold
0b00	—	FAULT (invalid)

This encoding is *not* two’s complement. The bit pattern 0b11 for zero is chosen so that the all-ones reset state of a register file initialises every register to hold—the semantically correct neutral value— without special reset logic. Negation maps to a simple bit swap: 0b01 $\leftrightarrow$ 0b10, leaving 0b11 invariant.

Proposition 1 (Negation as bit swap). *For encoding $\backslash(e(t)\backslash)$ as above, $\backslash(e(\mathrm{neg}\{t\})\backslash) = \{e(t)[0], e(t)[1]\backslash\}$ (swap bits).*

The Ternlang Language

Design principles

Ternlang is statically typed, expression-oriented, and compiled to BET bytecode. Three design decisions distinguish it from binary-native languages adapted to ternary:

Exhaustive three-way matching. Every match expression must cover all three trit arms. The compiler rejects non-exhaustive matches at parse time, eliminating an entire class of runtime error present in languages that bolt ternary logic onto binary pattern matching.

$\backslash(0\backslash)$ is active. The type system and semantics assign distinct meaning to $\backslash(-1\backslash)$ (conflict), $\backslash(0\backslash)$ (hold, active neutral), and $\backslash(+1\backslash)$ (truth). There is no null, no undefined, no NaN. A trit always carries a definite value.

Sparsity is a language feature. The @sparseskip directive marks a tensor operation as sparse-aware, routing the compiler to emit TSPARSE_MATMUL instead of TMATMUL. Sparsity is expressed in the source language, not discovered by the optimizer.

Core language constructs

```
fn classify(signal: trit) -> trit {
  match signal {
    -1 => conflict() // disagreement confirmed
    0  => hold()     // awaiting evidence
    +1 => truth()    // assertion confirmed
```

```

    }
}

// Route to TSPARSE_MATMUL at the ISA level
@sparseskip
let output: tritensor<8 x 8> =
    matmul(input, weights);

agent Voter {
    fn handle(msg: trit) -> trit {
        consensus(msg, hold())
    }
}

let v: agentref = spawn Voter;
send v truth();
let decision: trit = await v;

```

Type system

The core types are `trit` (a single balanced ternary value), `tritensor<N x M>` (an $N \times M$ matrix of trits stored on the tensor heap), and `agentref` (a handle to a running actor instance, local or remote). Struct types with trit/tensor fields are supported via field-name mangling in the register allocator.

Ternary Error Propagation

Ternlang introduces a postfix `?` operator for ternary-native error propagation, analogous to Rust's `?` operator but grounded in the three-valued semantics of the trit space. For any expression (e) of type `trit`, writing $(e\text{?})$ in a function body has the following semantics:

$$[e\text{?}] \triangleq \begin{cases} \text{return } -1 & \text{if } e = -1 \\ e & \text{otherwise} \end{cases} \text{label}\{\text{eq:propagate}\}$$

This elevates trit -1 (conflict) from a value to a structural signal: any function that receives a conflict can propagate it up the call chain without explicit match arms. At the BET level, `expr?` compiles to:

```

; expr already on stack: [val]
TDUP          ; [val, dup]
TJMP_NEG L1   ; consume dup; if -1 -> L1; else [val] continues
TJMP L2       ; skip early return
L1: TRET      ; return -1 to caller
L2: (continue with val on stack)

```

The parser disambiguates `?` from the ternary branching operator (also `?`) using two-token lookahead: `expr?` followed by `{` is an uncertain branch; `expr?` followed by any other token is a propagation operator.

```

fn validate_signal(x: trit) -> trit {
    let checked: trit = boundary_check(x)?; // propagate if conflict
    return range_check(checked)?;          // propagate again
}

fn main() -> trit {
    return validate_signal(1)?; // -1 if any check fails
}

```

Cross-File Module System

Ternlang's module system supports two resolution strategies, tried in order:

1. **Built-in stdlib** (compile-time embedded): `std::trit`, `std::math`, `std::tensor`, `std::io`, `ml::quantize`, `ml::inference` are embedded in the compiler binary via `include_str!`—zero filesystem I/O at runtime.
2. **User-defined modules**: a `ModuleResolver` walks use paths relative to the source file's directory. A declaration `use agents::voter;` resolves to `<source_dir>/agents/voter.tern`, parsed and merged into the program before semantic analysis.

Both strategies deduplicate: a module loaded by multiple use statements is parsed exactly once, and functions already present in the program are never duplicated. This makes `StdlibLoader::resolve()` idempotent and safe to call multiple times.

```

// agents/voter.tern
fn weighted_vote(a: trit, b: trit, c: trit) -> trit {
    consensus(consensus(a, b), c)
}

// main.tern
fn main() -> trit {
    use agents::voter;
    return weighted_vote(1, 0, -1);
}

```

Diagnostic Philosophy

Ternlang's error messages carry structured codes and ternary-philosophical commentary. Every diagnostic has a machine-readable code (for tooling and documentation lookup) and a human-readable nudge that frames the error in the language's formal architecture:

- [PARSE-001] Unexpected token — *"the lexer hit something it didn't expect. Check your syntax."*
- [PARSE-004] Non-exhaustive match — *"ternary has three states — cover all three or the compiler won't let you through."*
- [SCOPE-001] Undefined variable — *"hold state — declare before use."*

- [FN-002] Return type mismatch — *"ternary contracts are strict."*
- [BET-001] Stack underflow — *"you tried to pop a truth that wasn't there."*
- [PROP-001] ? on non-trit — *"only trit-returning functions can signal conflict. The third state requires a trit."*

This design ensures that newcomers encountering their first ternary error receive both actionable guidance and an introduction to the philosophy underpinning the type system.

The BET Instruction Set Architecture

Machine model

The BET VM is a stack-based machine with:

- **27 registers** (2 bits each), reset to 0b11 (hold)
- **Value stack**: unbounded, stores Value tagged union
- **Tensor heap**: indexed array of $(N \times M)$ trit matrices
- **Call stack**: return address stack for TCALL/TRET
- **Agent table**: maps type IDs to handler addresses and mailboxes
- **Carry register**: overflow from TADD

The choice of 27 registers reflects the ternary motif: $(27 = 3^3)$.

Instruction encoding

Instructions are variable-length, with a 1-byte opcode followed by 0–2 operand bytes. Table 3 lists the complete opcode map.

BET ISA opcode reference (selected)

Opcode	Mnemonic	Stack effect
0x00	THALT	—
0x01	TPUSH <i>t</i>	$(t \rightarrow t)$
0x02	TADD	$(a \ b \rightarrow s \ c)$
0x03	TMUL	$(a \ b \rightarrow t)$
0x04	TNEG	$(t \rightarrow \mathrm{neg}(t))$
0x05	TJMP_POS <i>addr</i>	$(t \rightarrow) \text{ (jmp if } (t=+1))$
0x06	TJMP_ZERO <i>addr</i>	$(t \rightarrow) \text{ (jmp if } (t=0))$
0x07	TJMP_NEG <i>addr</i>	$(t \rightarrow) \text{ (jmp if } (t=-1))$
0x08	TSTORE <i>r</i>	$(t \rightarrow) \text{ (reg}[(r)] \leftarrow t)$

Opcode	Mnemonic	Stack effect
0x09	TLOAD <i>r</i>	$\text{reg}[\text{r}]$
0x0b	TJMP <i>addr</i>	—
0x0e	TCONS	$\text{a} \ \text{b} \rightarrow \text{cons}(\text{a}, \text{b})$
0x0f	TALLOC <i>N M</i>	ref
0x10	TCALL <i>addr</i>	(push PC; jmp)
0x11	TRET	(pop PC; jmp)
0x20	TMATMUL	$\text{r}_A \ \text{r}_B \rightarrow \text{r}_C$
0x21	TSPARSE_MATMUL	$\text{r}_A \ \text{r}_B \rightarrow \text{r}_C$
0x22	TIDX	$\text{ref}[\text{i} \ \text{j} \rightarrow \text{t}]$
0x23	TSET	$\text{ref}[\text{i} \ \text{j} \ \text{t} \rightarrow]$
0x24	TSHAPE	$\text{ref}[\rightarrow \text{N} \ \text{M}]$
0x25	TSPARSITY	$\text{ref}[\rightarrow \text{count}]$
0x30	TSPAWN <i>id</i>	agentref
0x31	TSEND	$\text{agentref}[\text{m} \rightarrow]$
0x32	TAWAIT	$\text{agentref}[\rightarrow \text{t}]$

Sparse Ternary Inference

Ternary quantization

BitNet-style ternary weight quantization maps floating-point weights $w \in \mathbb{R}$ to $\hat{w} \in \{-1, 0, +1\}$ using a threshold τ :

$$\hat{w} = \begin{cases} +1 & w > \tau \\ 0 & |w| \leq \tau \\ -1 & w < -\tau \end{cases} \quad \tau = \frac{1}{2} \cdot \mathbb{E}[|w|] \quad \text{eq:quantize}$$

The resulting weight distribution is heavily concentrated at 0 (hold): typical language model weights at BitNet scale show 55%–65% zero elements after quantization with τ as defined in (eq:quantize).

The @sparseskip directive and TSPARSE_MATMUL

The key insight is that a multiply against a zero-weight trit produces zero regardless of the input value:

$$\text{mul}(\text{a}, 0) = 0 \quad \text{forall } \text{a} \in T \quad \text{eq:zero-mul}$$

In a dense $(N \times M)$ matrix multiply, every element contributes $(N \cdot M)$ multiplications. After ternary quantization with sparsity ρ (fraction of zero-weight elements), only $((1-\rho) \cdot N \cdot M)$ multiplications have non-trivial results.

TSPARSE_MATMUL (opcode 0x21) implements a sparse inner-product loop that tests the weight trit before executing the multiply:


```
for i in 0..N:
  for j in 0..M:
    for k in 0..K:
      w = W[k][j]
      if w == HOLD: continue // skip
      acc[i][j] += mul(A[i][k], w)
```

The directive @sparseskip in the source language is a compile-time annotation that causes the codegen to emit TSPARSE_MATMUL instead of TMATMUL for the following matmul() call. Sparsity awareness is thus a *source-language* property surfaced at the ISA level, not a runtime optimization that may or may not trigger.

Benchmark results

We evaluate sparse versus dense matrix multiply on (512×512) weight matrices quantized from normally distributed $\mathcal{N}(0,1)$ float weights using threshold ([eq:quantize](#)).

Sparse vs. dense ternary matmul (512×512)		
Metric	Dense	Sparse
Weight sparsity	0%	56.2%
Multiply operations	262,144	115,343
Skipped operations	0	146,801
Speedup	$(1.0 \times)$	$(\mathbf{2.27 \times})$

The $(2.27 \times)$ reduction in multiply operations is achieved without approximation: the result of TSPARSE_MATMUL is identical to TMATMUL by the identity ([eq:zero-mul](#)).

Real Hardware Validation: AVX2 SIMD Production Benchmarks

To validate the theoretical throughput analysis on physical silicon, we benchmarked ternary sparse matrix multiply against both a dense ternary baseline and an INT8 quantized baseline using AVX2 SIMD kernels deployed in a production Axum server. Throughput was measured in operations per second across three representative sparsity levels on a commodity x86-64 CPU.

AVX2 SIMD ternary inference — production hardware throughput		
Sparsity	vs. Dense Ternary	vs. INT8
10%	$(1.02 \times)$	—
50%	$(\mathbf{1.80 \times})$	—
90%	$(1.63 \times)$	$(\mathbf{3.0 \times})$

At 50% sparsity—typical of well-trained ternary models after L1 regularisation—the AVX2 kernel delivers a $\times 1.80$ throughput gain over dense execution. At 90% sparsity the advantage over INT8 quantization reaches $\times 3.0$, establishing that ternary is a strictly superior inference format at practical sparsity levels. These results are causal rather than correlational: perf-counter analysis confirms that the throughput gain tracks the zero-weight skip rate, and statistical validation confirms the measurement is stable across repeated runs.

The benchmark establishes a hardware-verified empirical foundation for the central claim of the @sparseskip design: surfacing sparsity at the ISA level produces measurable throughput gains on commodity hardware without approximate arithmetic, hardware reconfiguration, or GPU acceleration.

Hardware Backend

Verilog-2001 primitives

The ternlang-hdl crate generates synthesisable Verilog-2001 modules from the BET ISA. Each trit becomes a $[1:0]$ bus. The core primitives are:

- trit_neg: bit swap (assign $y = \{a[0], a[1]\}$)
- trit_cons: assign $y = (a == b) ? a : 2'b11$
- trit_mul: zero-skip detect with combinational multiply
- trit_add: 9-case Verilog case with carry
- trit_reg: synchronous D-register with reset-to-hold
- bet_alu: top-level ALU selecting among primitives by op code

Sparse matmul array

The synthesisable sparse matmul array instantiates an $(N \times N)$ grid of processing elements. Each cell contains a weight register and a clock-gate signal:

```
wire [1:0] w_ij = weight_reg[i][j];
wire skip = (w_ij == 2'b11); // hold = zero weight
wire [1:0] contrib = skip ?
    2'b11 : // hold if skipped
    trit_mul(a_i, w_ij);
```

Clock-gating on the skip signal prevents switching activity in zero-weight cells, delivering dynamic power reduction proportional to sparsity.

BET processor

The full `bet_processor` module wires together:

- `bet_regfile`: 27-register file ((27×2) bits)
- `bet_pc`: 16-bit program counter with load port for jumps
- `bet_control`: single-cycle decode, all 51 opcodes mapped to control signals (`alu_op`, `reg_we`, `mem_re`, `mem_we`, `pc_load`, `is_call`, `is_ret`, `is_halt`, `is_spawn`, `is_send`, `is_await`)

An Icarus Verilog simulation wrapper (`BetSimEmitter`) generates complete self-contained testbenches from BET bytecode, including inline ROM initialization, reset sequencing, and VCD waveform export for GTKWave.

Native Compilation: AST-to-C Backend

The `ternlang-codegen` crate provides a second compilation target alongside BET bytecode: a self-contained C11 source file that can be compiled with any standard C compiler. This opens three new deployment paths: native-speed execution, cross-compilation to embedded targets, and human-readable output for inspection and security audit.

Transpilation Architecture

The `CTranspiler` operates directly on the abstract syntax tree (Program) produced by the parser, bypassing the BET VM entirely:

```
[\texttt{.tern} \xrightarrow{\text{parse}} \text{AST} \xrightarrow{\text{CTranspiler}} \text{.c} \xrightarrow{\text{gcc/clang}} \text{native binary} \label{eq:ctranspile}]
```

The generated C file is fully self-contained: a header of inline trit primitives using `int8_t` with values $\{-1, 0, +1\}$ is prepended, covering `trit_add`, `trit_mul`, `trit_neg`, `trit_consensus`, `trit_abs`, and the built-in constants `trit_truth()`, `trit_hold()`, `trit_conflict()`.

Language Construct Mapping

The transpiler maps every ternlang construct to its natural C equivalent:

- `match` $\rightarrow$ `switch (int)` with case `-1`, case `0`, case `1` arms
- `struct` definitions $\rightarrow$ `typedef struct { ...}`
- `fn` $\rightarrow$ C function with forward declaration (enables mutual recursion without re-ordering)

- `loop / while / for` $\rightarrow$ `for(;;), while(1)` with inner ternary condition dispatch
- `expr?` $\rightarrow$ `__TERN_PROPAGATE(expr)` macro slot, enabling the caller to define the early-return idiom in C
- `main()` $\rightarrow$ `tern_main()` (to avoid clashing with C's entry point); a generated `main()` calls `tern_main()` and translates the trit result to an exit code

```
// ternlang source
fn decide(x: trit) -> trit {
  match x {
    -1 => conflict()
    0  => hold()
    +1 => truth()
  }
}

// generated C
trit decide(trit x) {
  switch ((int)x) {
    case -1: return trit_conflict(); break;
    case 0:  return trit_hold();      break;
    case 1:  return trit_truth();     break;
  }
}
```

The C backend does not yet support the actor model (`spawn/send/await`) or tensor heap operations, which remain BET-VM specific. These emit comments in the output, making the generated file auditable even for unsupported constructs.

Actor Model and Distributed Runtime

Local actors

Ternlang implements the actor model via three ISA primitives: `TSPAWN` creates an agent instance and returns an `agentref`; `TSEND` enqueues a trit message in the agent's mailbox; `TAWAIT` dequeues a message, invokes the handler, and returns the trit result. The mailbox is a FIFO (`VecDeque<Value>`) allocated per agent instance.

Distributed actors

The `ternlang-runtime` crate extends the actor model to multi-node systems via TCP transport. A `TernNode` binds a TCP port, maintains a peer connection map, and exposes `remote_send/remote_await` over a newline-delimited JSON wire protocol:

```
let remote_voter: agentref =
  spawn remote "192.168.1.42:7373" Voter;
```

```
send remote_voter truth();
let r: trit = await remote_voter;
```

The wire protocol uses four message types (Send, Await, Reply, Error) serialised as tagged JSON objects, enabling interoperability with non-Rust implementations.

Ternary AI Reasoning Toolkit

Modern AI agents are structurally binary in their decision loops: evidence is collapsed to a scalar, thresholded, and converted to a Boolean act/wait flag. The three-valued nature of evidence—affirm, hold, reject—is treated as a side-effect of confidence scores rather than a first-class type.

Phase 8 of the TIS introduces five primitives in `ternlang-ml` that make the decision loop architecturally ternary.

Deliberation Engine

The `DeliberationEngine` models iterative evidence accumulation using an exponential moving average (EMA):

$$S_r = \alpha \cdot e_r + (1 - \alpha) \cdot S_{r-1}, \quad \alpha \in (0, 1]$$

$\text{\label{eq:ema}}$

where (e_r) is the mean of new evidence signals in round (r) . The engine commits to a ternary decision only when the confidence of (S_r) exceeds a target threshold; otherwise it remains in the hold state and requests a further evidence round. This models the human “let me think about it” behaviour—a native ternary computation rather than a binary timeout.

Coalition Vote

`coalition_vote()` aggregates (N) independent agent verdicts—each a trit with an associated confidence and weight—into a single result with quorum, dissent, and abstain statistics:

$$\hat{t} = \text{sign}\left(\frac{\sum_i w_i \cdot c_i \cdot t_i}{\sum_i w_i \cdot c_i}\right) \quad \text{\label{eq:coalition}}$$

where $(t_i \in \{-1, 0, +1\})$, $(c_i \in [0, 1])$ is confidence, and $(w_i > 0)$ is importance weight.

Action Gate

The action gate enforces a structural separation between hard constraints and soft preferences. Each `GateDimension` carries an evidence signal, a weight, and an optional `hard_block` flag. A hard-block dimension with negative evidence vetoes the action

unconditionally, regardless of all other dimensions:

$$\text{verdict} = \begin{cases} \text{Block} & \exists d_i : \text{hard_block}(d_i) \\ \text{trust}_{d_i} = -1 & \text{sign}(\bar{S}) \text{ and otherwise} \end{cases}$$

$$\text{gate} \text{ label{eq:gate}}$$

This is the structural analogue of a safety interlock: the veto is unconditional, not probabilistic.

Scalar Temperature and Hallucination Score

`scalar_temperature()` bridges the ternary decision to the LLM sampling temperature domain: affirm at high confidence maps to a low temperature (focused generation), hold maps to a mid temperature (exploratory generation), and reject maps to near-zero (cautious refusal).

`hallucination_score()` maps the variance of an evidence signal vector to a trust trit: low variance (consistent signals) yields trust $(+1)$; high variance yields (-1) , signalling that the agent's signals are incoherent and should not be acted upon.

MoE-13 Ternary Orchestrator

The MoE-13 architecture is implemented in the `ternlang-moe` crate. It realises a ternary Mixture-of-Experts head whose routing, synthesis, safety gate, and memory are all expressed natively in balanced ternary semantics.

Six-Dimensional Competence Space

Each expert is characterised by a *competence vector* $\mathbf{v} \in [-1,1]^6$ across six axes:

$$\mathbf{v} = [\text{syntax}, \text{world}, \text{reason}, \text{tool}, \text{persona}, \text{safety}] \text{ label{eq:competence}}$$

The synergy between two experts is defined as the complement of their cosine similarity—orthogonal experts are maximally complementary:

$$\sigma(\mathbf{v}_i, \mathbf{v}_j) = \frac{1 - \cos(\mathbf{v}_i, \mathbf{v}_j)}{2} \in [0, 1] \text{ label{eq:synergy}}$$

Dual-Key Synergistic Routing

For each candidate expert pair $((i, j))$, the routing score combines relevance to the query vector $\mathbf{q}$ with inter-expert synergy:

$$\text{score}(i, j) = \rho_i(\mathbf{q}) \cdot \rho_j(\mathbf{q}) \cdot \sigma(\mathbf{v}_i, \mathbf{v}_j) \text{ label{eq:routing}}$$

where $\rho_k(\mathbf{q}) = \max(0, \cos(\mathbf{v}_k, \mathbf{q}))$ is the relevance of expert k to query $\mathbf{q}$. The pair with the highest score is selected. Crucially, two highly relevant but similar experts are penalised by low σ ; the router favours complementary coverage over redundant strength.

$(1+1=3)$ Triad Synthesis

After expert evaluation, an emergent *triad field* is synthesised from the two selected competence vectors and the routing synergy:

$$\mathbf{E}_k = \sigma_{ij} \cdot \frac{\mathbf{v}_i + \mathbf{v}_j}{2}$$

$\text{\label{eq:triad}}$

This is the formal expression of $(1+1=3)$: two expert signals at high synergy produce a third signal—the emergent field $\mathbf{E}_k$ —that neither expert could produce in isolation. The field modulates the subsequent vote and the LLM temperature hint.

Safety Hard Gate and Axis-6 Veto

Dimension 6 of every competence vector is the *safety axis*. Before any vote is computed, the safety projection of the triad field is evaluated:

$$\text{veto} \iff E_{k,\text{safety}} < \theta_s \quad \text{\label{eq:veto}}$$

where θ_s is the configurable safety threshold (default -0.3). A veto immediately returns trit -1 at confidence 1.0 and logs the event to the Axis-tier memory for audit. No downstream reasoning can override it.

Hold State and Tiebreaker

When the weighted vote yields trit 0 or aggregate confidence falls below a hold threshold, the orchestrator does not immediately return. It invokes a tiebreaker expert—selected by highest reasoning-axis score among inactive experts—and re-votes with up to $\max_{\text{active}}=4$ experts total. Only if the tiebreaker also fails to resolve does the orchestrator return the hold state to the caller, signalling that further evidence is needed.

Three-Tier Memory Mesh

The orchestrator maintains three memory tiers:

- **Node** (TTL: seconds): LRU-bounded volatile store for within-session context. Evicts on capacity overflow.
- **Cluster** (TTL: minutes): routing frequency counters enabling mode-collapse detection ($\text{risk} = \max_{\text{pair}} / \text{total}$).
- **Axis** (persistent): immutable audit log of safety vetoes with timestamp, expert identity, and query hash. Global priors over expert performance.

Nine-Step Inference Pipeline

The full orchestration pass executes in nine deterministic steps: (1) encode query to evidence vector; (2) dual-key route to best expert pair; (3) evaluate both experts independently; (4) synthesise triad field via ([\[eq:triad\]](#)); (5) safety hard gate via ([\[eq:veto\]](#)); (6) weighted trit vote with synergy amplification; (7) hold detection; (8) optional tiebreaker invocation; (9) return `OrchestrationResult` and update all three memory tiers.

Thirteen-Agent Dual-Signal Deliberation

The `AgentHarness` in `ternlang-moe/src/agents/` extends the 9-step orchestration pass with a structured deliberation layer of 13 specialised agents, each computing two independent signals:

$$\begin{cases} +1 & \text{if positive signal} \geq \theta^+ \text{ and negative signal} < \theta^- \\ -1 & \text{if negative signal} \geq \theta^- \text{ and positive signal} < \theta^+ \\ 0 & \text{(stasis: signals in equilibrium)} \end{cases}$$

[\[eq:dual_signal\]](#)

where θ^+ and θ^- are agent-specific thresholds. The 13 agents cover: **Syntax** (structural token analysis, bracket balance), **WorldKnowledge** (domain vocabulary density, proper noun frequency), **DeductiveReason** (premise + conclusion marker co-occurrence), **InductiveReason** (example density, generalisation leap detection), **ToolUse** (imperative verb strength, passive construction penalty), **Persona** (depersonalisation detection), **Safety** (hard risk keywords $\text{trit}(-1)$ at confidence 0.99), **FactCheck** (overconfidence markers $\text{trit}(\text{to})$ hallucination flag), **CausalReason** (requires both cause and effect markers), **AmbiguityRes** (hard-vague ≥ 2) markers $\text{trit}(-1)$, **MathReason** (impossible operations e.g. divide-by-zero $\text{trit}(-1)$), **ContextMem** (fresh-start signals vs. anaphoric reference density), and **MetaSafety** (injection pattern detection at confidence 0.99).

The dual-signal design is philosophically significant: $\text{trit}(0)$ is not a default or a fallback from a failed comparison. It is *active stasis*—the system’s honest declaration that positive and negative evidence are balanced and further information is required before committing.

Introspective Hold: The Stable Attractor

The `run_introspective()` method on `AgentHarness` implements a *stable attractor* for the hold state. Once reached, a stable hold is permanent within the deliberation turn—it cannot be overridden by additional tiebreaker invocations. The stable attractor condition is:

$$\text{stable_hold} \iff \text{affirm_count} - \text{conflict_count} \leq 1 \wedge \text{engaged} \geq 4$$

[\[eq:stable_hold\]](#)

This formalises the philosophical claim: when at least four independent agents have deliberated and the signal balance between affirmation and conflict is indistinguishable, the appropriate epistemic response is to hold—not to choose arbitrarily between $\text{trit}(+1)$

and $\neg$. The hold state is returned with `is_stable_hold: true` and a human-readable `hold_reason` string.

The aggregate verdict structure carries full deliberation metadata:

```
pub struct AggregateVerdict {
    pub trit: i8,
    pub confidence: f32,
    pub verdicts: Vec<ExpertVerdict>,
    pub is_stable_hold: bool,
    pub hold_reason: Option<String>,
    pub affirm_count: usize,
    pub conflict_count: usize,
    pub hold_count: usize,
}
```

Safety and meta-safety form a hard gate: the evidence vector produced by `to_evidence_vector()` maps axis 6 to $\min(\text{Safety.confidence}, \text{MetaSafety.confidence})$, ensuring the safety dimension is always the most conservative signal in the system.

Orchestrate-Full: Thirteen-Substage Pipeline

The `orchestrate_full()` method on `TernMoeOrchestrator` composes the 13-agent deliberation with the 9-step routing pipeline into a unified end-to-end pass:

1. Run all 13 agents via `AgentHarness::run_introspective()`
2. Check for stable attractor hold — return immediately if true
3. Map 13 verdicts to 6D evidence vector via `to_evidence_vector()`
4. Check safety hard gate on safety axis before MoE routing
5. Run standard 9-step MoE orchestration with enriched evidence
6. Return `OrchestrationResult` with stable hold flag propagated

This creates a two-tier system: the fast ternary language models (13 agents, keyword-level deliberation) form a pre-filter that enriches the evidence before the slower but more semantically powerful MoE routing pass. Most queries that are clearly safe and unambiguous complete in the agent tier; only complex or borderline queries reach the full routing pipeline.

MCP Exposure

The orchestrator is exposed as three MCP tool calls: `moe_orchestrate` runs a full pass and returns the complete result including routing pair, triad field, verdicts, temperature, and prompt hint; `moe_deliberate` wraps the EMA deliberation engine for multi-round iterative reasoning; `trit_action_gate` exposes the hard-block gate as a standalone safety check. Any MCP-compatible AI client connected to `ternlang-mcp` gains access to the full ternary reasoning stack as tool calls.

Albert MoE-13: Autonomous Training on Consumer CPU Hardware

Albert MoE-13 v2.0.0 is the reference implementation of the ideas developed in Sections 5–10: a fully ternary Mixture-of-Experts language model trained end-to-end without GPU hardware. It serves as an empirical existence proof that the ternary execution substrate enables practical neural training and inference on commodity consumer hardware.

Architecture

Albert MoE-13 v2.0.0 architecture

Parameter	Value
Hidden size	256
Experts per layer	12
Active experts per token	3 (Top-3 sparse)
Depth	3 layers (3L training phase)
Total parameters	$\approx 22\text{M}$
Active parameters per token	7.6M (34.5%)
Weight format	ternary $\{-1, 0, +1\}$ throughout
Checkpoint size (float32 on disk)	91 MB
Packed ternary deployment estimate	$\approx 6\text{ MB}$

All 22 million parameters are stored as discrete values $w \in \{-1, 0, +1\}$. Top-3 sparse routing over 12 experts means 7.6M parameters are active per forward pass—a 65.5% parameter reduction per inference step, realised by the same TSPARSE_MATMUL mechanism described in Section 5.

The EvolutionManager: Autonomous Architectural Growth

The most significant new contribution of v2.0.0 is the EvolutionManager: an autonomous state machine that monitors training dynamics and triggers *in-situ* architectural expansion without operator intervention. The protocol executes in four phases:

1. **Monitor:** track epoch-over-epoch loss delta against configurable plateau and mastery thresholds
2. **Trigger:** when the delta satisfies both conditions simultaneously, enqueue a surgery event

3. **Surgery:** clone the deepest transformer layer into a new layer $(L+1)$ via Net2Net safe-copy initialisation—an identity-preserving weight transformation that produces no loss spike at the architectural boundary
4. **Resume:** training continues from the expanded architecture; all accumulated linguistic structure from prior epochs is preserved in the cloned weights

On 2026-05-06, the EvolutionManager triggered the first autonomous surgery: a $(3L \rightarrow 5L)$ depth expansion. The checkpoint grew from 41 MB to 91 MB. Loss at the transition boundary showed no regression. This represents the first documented instance of a fully ternary neural model expanding its own architecture without human intervention—the substrate evolving itself.

Training Efficiency: The 50 $\times$ Speedup

During v2.0.0 development, a critical bug was identified in the gradient accumulation loop: `backward_step()` was called once per *micro-batch* (16 times per logged batch) rather than once per full accumulation cycle, producing 16 separate under-accumulated optimizer updates per batch. Correcting this produced:

- **Wall-clock:** batch time $(\approx 50\text{s} \rightarrow \approx 2\text{s})$ ($(25\times)$ per-batch improvement)
- **Loss:** inflated sum (≈ 112) true cross-entropy $(\approx 6.5) \rightarrow (7.0)$
- **Gradient quality:** smooth descent from Epoch 1, consistent low-loss breakouts from Epoch 5

The combined effect on useful optimiser throughput is approximately $(50\times)$.

Ternary Training Innovations

Three training-time optimisations reduce compute overhead without degrading model quality:

L1 Sparsity Regularisation ($(\lambda = 10^{-5})$): added to the training loss, this term encourages the model to learn *where* to be sparse during training rather than post-hoc, producing higher signal concentration in active expert weights.

Per-Layer Threshold Gradient: Layer 0 uses quantisation threshold $(\tau_0 = 0.01)$ (dense, learning syntax and surface structure); the deepest layer uses $(\tau_d \geq 0.03)$ (sparse, learning higher-order abstractions). Intermediate layers interpolate linearly, mirroring the known depth-specialisation structure of biological neural circuits.

Gamma Cache in TernaryLinear: the mean activation $(\mu = \mathrm{mean}_{all}(W))$, previously recomputed on every forward call, is now cached and refreshed every 20 forward calls. This eliminates $(36\times)$ redundant full-weight reductions per batch at 256H. The causal attention mask is cached in `Attention` and rebuilt only on sequence-length change.

Training Results

Albert MoE-13 v2.0.0 — 16-epoch training summary

Metric	Value
Global epochs completed	16
Total batches logged	4,864
Epoch 1 average loss	8.3784
Epoch 15 average loss	6.5213
All-time best (single batch)	2.1353 (Epoch 9)
Effective perplexity at best batch	≈ 8.5
Low-loss breakouts (< 5.0) Epoch 1–3	0
Low-loss breakouts (< 5.0) Epoch 15	6
Average epoch duration	12–13 min
Training hardware	HP ZBook 2013 $\cdot$ Intel CPU $\cdot$ no GPU

The all-time best of 2.1353 (Epoch 9, single batch) corresponds to an effective perplexity of ≈ 8.5 on that batch—achieved when a specialised expert’s routing aligned precisely with a repetitive passage from the King James Bible corpus. This demonstrates that individual MoE experts have reached genuine pattern mastery on specific token distributions; the trit is not approximating a distribution, it is encoding a fact.

The loss variance pattern is characteristic of healthy MoE training: variance expands as experts differentiate (Epochs 4–10, $\sigma \approx 0.54$) then partially compresses as routing stabilises (Epoch 11, $\sigma = 0.45$). The floor of achievable loss per batch has monotonically decreased across all 16 epochs.

The training hardware is intentionally unremarkable: a 2013 HP ZBook, an Intel CPU, and no GPU. This is a design goal, not a constraint. If a 22M-parameter ternary MoE model can be trained to expert-level pattern mastery on a decade-old consumer CPU, the case for ternary as a democratising inference substrate is empirical rather than theoretical.

Last Look Back (LLB) Safety Protocol

The Last Look Back (LLB) protocol, implemented as a standalone MCP server (albert-llb-mcp), is a deterministic, policy-based filesystem safety gate that evaluates write operations against a configurable rule set before execution. LLB is stateless, deterministic (identical input always yields identical verdict), and ternary-native: its verdict is $\{ \text{trit} \in \{-1, 0, +1\} \}$ (reject / hold / affirm). It operates independently of the model it protects, ensuring that safety gates in sovereign AI systems remain auditable, reproducible, and immune to model-level compromises.

Qutrit Neural Networks and Ternlang

The Qutrit Neural Network (QNN) framework establishes a formal correspondence between quantum qutrit states and balanced ternary values. The three qutrit basis states $\lvert -\rangle$ (decay), $\lvert 0 \rangle$ (stasis), and $\lvert + \rangle$ (growth) map exactly onto the ternlang trit space:

$$\lvert - \rangle \leftrightarrow -1, \quad \lvert 0 \rangle \leftrightarrow 0, \quad \lvert + \rangle \leftrightarrow +1 \quad \text{label}\{\text{eq:qnn_map}\}$$

This correspondence is not coincidental. Balanced ternary's symmetric structure around zero and its active neutral state mirror the physical properties of three-state quantum systems, making ternlang a natural implementation substrate for QNN inference.

Tesseract Recursive Syntax Stabilisation

The QNN paper introduces *Tesseract Recursive Syntax* (TRS), a stabilisation mechanism that prevents qutrit networks from collapsing into binary modes under iterative refinement. In ternlang, TRS corresponds to the introspective hold architecture (Section 10.9): when affirmation and conflict signals are balanced, the system stabilises at trit $\lvert 0 \rangle$ rather than forcing a premature commitment.

QNN Example Suite

The ternlang repository contains 15 QNN-inspired .tern example programs (files 251–265) covering:

- Qutrit superposition encoding and collapse simulation
- Ternary attention mechanisms with $\{-1, 0, +1\}$ key-query products
- Qutrit activation functions replacing ReLU/GELU in ternary networks
- Phase-encoding of temporal signals as trit sequences
- TRS stabilisation loops using ternlang's loop construct and the ? propagation operator for collapse detection
- Qutrit gradient approximation via balanced ternary finite differences

These examples demonstrate that ternlang's syntax is sufficient to express QNN computations natively, without any binary adaptor layer. The @sparseskip directive applies directly to qutrit weight matrices: at realistic qutrit sparsity levels ($\approx 33\%$ zero states at initialisation), the TSPARSE_MATMUL kernel provides the same $8\times$ speedup over dense computation observed in classical BitNet workloads.

Related Work

Balanced ternary computing. Knuth provides the mathematical foundation. The Setun computer (1958) demonstrated physical ternary hardware. Modern revivals include academic hardware studies at the University of South-Eastern Norway, which targets C-

to-ternary compilation for EDA tools—an approach that inherits binary-native semantics and misses the symmetric properties of balanced ternary that ternlang is designed to exploit natively.

Ternary emulators. Several open-source projects implement ternary VMs in Python or JavaScript (Brandon Smith’s 9-trit RISC simulator; Owlet, an S-expression ternary interpreter in Node.js). These provide single-layer implementations without compiler, ML kernels, or hardware backends. ternlang-compat provides assembler-level bridges to both.

Ternary quantization for neural networks. BitNet and subsequent work (e.g., BitNet b1.58) demonstrate that large language models can be trained with ternary weights while retaining competitive perplexity. The present work is the first to surface this property as a first-class ISA opcode (TSPARSE_MATMUL) rather than a software library optimization.

Quantum ternary. Qutrits (3-level quantum systems) map naturally to trit values $\{-1, 0, +1\}$. The BET encoding and tritensor type system are structurally compatible with qutrit state spaces; we leave the formal mapping to future work.

The Ternary Computing Ecosystem

A stated goal of ternlang is to serve as a *convergence point* for the fragmented ternary computing field. Table 8 maps active projects to their interoperability path via ternlang-compat.

Ecosystem compatibility map		
Project	Interface	Bridge
Brandon Smith 9-trit sim	.tasm assembly	TasmAssembler $\rightarrow$ BET bytecode
Owlet (S-expr)	S-expression syntax	OwletParser $\rightarrow$ ternlang AST
USN / EDA tools	C-to-ternary	Academic whitepaper; ISA interop
Trit-Rust crate	Rust i8 trits	Superseded by ternlang-core
Q-Ternary	Qutrits	tritensor state model (future)

Beyond runtime interoperability, ternlang aims to serve as a *technical archive* for prior ternary computing work. Brandon Smith’s Python 9-trit RISC simulator demonstrated a complete fetch-decode-execute cycle in balanced ternary—an existence proof that received little downstream adoption due to the absence of a compiler ecosystem. The Owlet interpreter proved that S-expression evaluation maps cleanly to a three-valued substrate. Both contributions are honoured in the examples/ corpus: 09_risc_fetch_decode.tern translates Smith’s pipeline decision logic into native BET language syntax, and 13_owlet_bridge.tern models cooperative ternary evaluation with hold-as-suspension semantics. In this way, prior work that was “slowly forgotten” acquires a working executable form—runnable on the same VM that drives the MoE-13 orchestrator.

Implementation Status

Ternlang is implemented in Rust and comprises the following crates:

Workspace crates, test counts, and descriptions (v2.0.0)

Crate	Tests	Description
ternlang-core	43	Lexer, parser, AST, semantic checker, BET VM (51 opcodes)
ternlang-ml	21	BitNet quantization, sparse/CSC matmul, AVX2 SIMD kernels, deliberation engine, coalition vote, hallucination score
ternlang-moe	24	MoE-13 orchestrator: dual-key routing, triad synthesis, 3-tier memory, 13 dual-signal agents, introspective hold
ternlang-codegen	8	AST-to-C11 transpiler; match/struct/propagation
ternlang-test	10	Full-pipeline test framework
ternlang-hdl	34	Verilog-2001 codegen, BET RTL simulator
ternlang-runtime	4	Distributed TCP actor runtime (TernNode)
ternlang-compat	29	.tasm assembler (9-trit RISC), Owlet S-expr parser
ternpkg	5	Package manager, GitHub-backed registry
ternlang-lsp	—	LSP 3.17: hover, completion (19 snippets), diagnostics
ternlang-mcp	—	MCP server, 34 free tools , live on Smithery
ternlang-api	—	REST API (18 endpoints), KPI pipeline, SSE streaming
albert-llb-mcp	—	LLB safety gate: deterministic ternary verdict, stateless
albert-cli	—	Full Rust TUI agent CLI (ratatui); multi-provider; voice STT
exatern	—	Phase 11.6: SIMD trit packing, zero-copy array views

The CI-verified test suite comprises **138 fast tests across 5 self-contained crates** (core: 43, codegen: 8, moe: 24, compat: 29, hdl: 34), all passing. All crates are published to `crates.io` under open-core licensing (LGPL-3.0-or-later for the compiler and CLI; BSL-1.1 converting to LGPL in 2030 for the ML engine and MoE orchestrator). The live API is deployed at `https://ternlang.com/mcp` (Fly.io, Frankfurt) serving both the marketing website and the MCP endpoint from a single Rust binary.

Phase 17 — BET VM in WebAssembly. The BET VM has been compiled to WebAssembly and executes live in the browser via TernStudio. Users can write, compile, and run `.tern` programs without local installation. The WASM build is embedded at

compile time via `include_str!` and served directly. This is the first public demonstration that the complete ternary execution pipeline runs client-side on standard web hardware.

TernStudio. The browser IDE received two foundational additions in v2.0.0: the *Ternary Actuator Protocol* (TAP), an autonomous tool-call interception layer that suspends State-0 (uncertainty) nodes until a human or upstream agent resolves the signal; and *Liquid Time*, a DAW-style multiverse timeline scrubber enabling backward-scrub through simulation history without state corruption.

The open-core library now spans **28,500+ executable .tern programs**, making TIS the largest public ternary programming library in existence. The public website (ternlang.com) is fully bilingual EN/DE. The VS Code extension (`ternlang-0.1.0`) is published on the Open VSX Registry under publisher `rfl-irfos`. Continuous integration via GitHub Actions automatically builds, tests, and deploys to Fly.io on every push to main.

Conclusion and Future Work

We have presented Ternlang v2.0.0—*The Autonomous Substrate*—a full-stack balanced ternary execution architecture spanning language design, ISA, virtual machine, hardware backend, distributed runtime, native C compilation, AI reasoning infrastructure, and autonomous neural training. Six layers of contribution are unified under the same foundational primitive—the trit $\{ -1, 0, +1 \}$ with an active neutral state:

1. **Language completeness:** exhaustive three-way matching, the $\oplus$ error propagation operator, a real cross-file module system (stdlib built-in + filesystem-relative user modules), structured diagnostic codes, and all major control flow constructs—making ternlang a viable general-purpose programming language rather than a research prototype.
2. **Execution substrate:** TSPARSE_MATMUL achieves a $(2.27\times)$ reduction in multiply operations at baseline sparsity; a CSC kernel with Rayon reaches $(122\times)$ at high sparsity (512² matrix, 99% sparsity, release mode). A native C11 compilation path via `ternlang-codegen` provides an alternative to VM interpretation.
3. **Reasoning primitives:** the five-tool Ternary AI Reasoning Toolkit (DeliberationEngine, coalition vote, action gate, scalar temperature bridge, hallucination score) makes AI agent decision loops structurally three-valued.
4. **MoE orchestration:** the MoE-13 orchestrator extends with 13-agent dual-signal deliberation, an introspective hold (stable attractor when $(|\text{affirm}| - |\text{conflict}|) \leq 1$) across (≥ 4) agents), and an `orchestrate_full()` pipeline that pre-filters queries through the agent tier before MoE routing.
5. **Deployment and tooling:** 34 free MCP tools live at <https://ternlang.com/mcp>; the BET VM runs in the browser via WebAssembly; a VS Code extension with LSP, formatter, and REPL; and 138 CI-verified tests across 5 fast crates (28,500+ .tern stdlib programs) validate the full stack.

6. **Autonomous training:** Albert MoE-13 v2.0.0 (22M ternary parameters, Top-3/12 MoE, 256H) achieved an all-time best cross-entropy of **2.1353** (perplexity ≈ 8.5) after 16 global epochs of autonomous training—including one self-triggered $(3L \rightarrow 5L)$ depth expansion via the EvolutionManager—running entirely on a 2013 consumer CPU with no GPU. Hardware-verified AVX2 SIMD benchmarks confirm $(1.80\times)$ throughput at 50% sparsity and $(3.0\times)$ over INT8 at 90% sparsity.

The title of this paper carries a provocation: *The Death of the Bit*. It is not an obituary for the trit—it is an obituary for the bit. It is a claim about what happens when ternary computation matures: the binary bit, the foundational primitive of five decades of computing, stops being the only option and becomes a special case of a more expressive substrate. The goal of the Ternary Intelligence Stack is to make that invisibility possible—to build the substrate so completely that the three-valued logic becomes the ground truth, not the exception. We believe Europe’s path to technological sovereignty runs through this kind of foundational work: not chasing the scaling frontier, but owning the substrate beneath it.

The VS Code extension (ternlang-0.1.0) has been published on the Open VSX Registry (publisher rfi-irfos), making ternlang tooling available to VS Code, VSCodium, Gitpod, and any editor backed by the Open VSX API. The SPRIND Next Frontier AI Challenge submission (deadline 2026-05-15) documents the full TIS architecture for European public funding consideration. Near-term targets include submission to the VS Code Marketplace and academic outreach to the USN group (Bos & Gundersen) for joint hardware-software co-design on the memristor backend.

Further research directions include:

- **Type inference:** Hindley-Milner style inference to eliminate explicit type annotations, reducing boilerplate for trit-returning functions
- **FPGA synthesis:** full bet_processor targeting Xilinx Artix-7 and Lattice ECP5, benchmarked against the pure-Rust RTL simulator
- **End-to-end training:** native ternlang training loop with BitNet-style gradient quantization, ternary weight persistence, and backward pass in balanced ternary arithmetic
- **QNN acceleration:** formal integration of the Qutrit Neural Network framework with TSPARSE_MATMUL as the inner-loop kernel, targeting quantum-adjacent hardware via the TRS stabilisation pattern
- **Memristor backend:** integration with physical ternary storage via the USN uMemristorToolbox
- **Community bootstrap:** coordinated engagement with the broader ternary computing ecosystem (Brandon Smith 9-trit, Owlet S-expr, Trit-Rust, BitNet b1.58 community)

The ternary computing field has been fragmented for decades. Ternlang is designed to be the substrate where those fragments converge—from ISA to orchestrator, from single trit to 13-expert MoE, from quantum-adjacent qutrit networks to physical memristor hardware. The philosophy is simple: binary sees yes and no; ternary also sees *not yet*—and that changes everything.

1. RFI-IRFOS — Research Focus Institute, Elisabethnergasse 25, 8020 Graz, Austria.
Patent Pending A50296/2026. Repository: <https://github.com/eriirfos-eng/ternary-intelligence-stack>